


PAINFULLY SLOW: The Offline Least Common Ancestor Problem

Matthew Kilgore

May 5, 2016

Example



Let's get some basic facts out of the way. Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$:

Let's get some basic facts out of the way. Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$:

- We say y is a **descendant** of x if the unique path

$$P = (x = x_0, x_1, \dots, x_n = y)$$

from x to y is such that x_i is x_{i+1} 's parent.

Let's get some basic facts out of the way. Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$:

- We say y is a **descendant** of x if the unique path

$$P = (x = x_0, x_1, \dots, x_n = y)$$

from x to y is such that x_i is x_{i+1} 's parent.

- The **least common ancestor** of x and y is the vertex v furthest from the root such that x and y are descendants of v .

Let's get some basic facts out of the way. Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$:

- We say y is a **descendant** of x if the unique path

$$P = (x = x_0, x_1, \dots, x_n = y)$$

from x to y is such that x_i is x_{i+1} 's parent.

- The **least common ancestor** of x and y is the vertex v furthest from the root such that x and y are descendants of v .
- Alternatively, the least common ancestor of x and y is the vertex v on the unique path from x to y closest to the root.

The Question

Least Common Ancestor Problem

Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$, efficiently compute the least common ancestor of x and y .

The Question

Least Common Ancestor Problem

Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$, efficiently compute the least common ancestor of x and y .

- This arises in applications such as when examining a class inheritance hierarchy.

Least Common Ancestor Problem

Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$, efficiently compute the least common ancestor of x and y .

- This arises in applications such as when examining a class inheritance hierarchy.
- Sometimes it makes sense to find the pairs of vertices to query before actually querying.

The Question

Least Common Ancestor Problem

Given a rooted tree $T = (V, E)$ and vertices $x \neq y \in V$, efficiently compute the least common ancestor of x and y .

- This arises in applications such as when examining a class inheritance hierarchy.
- Sometimes it makes sense to find the pairs of vertices to query before actually querying.

Offline Least Common Ancestor Problem

Given a rooted tree $T = (V, E)$ and m pairs of vertices $x \neq y \in V$, efficiently compute the least common ancestor for each of the m pairs x and y .

[PAUSE FRAME]

Who?



R. E. Tarjan

Who?



R. E. Tarjan



H. N. Gabow

Theorem

We can solve the Offline Least Common Ancestor Problem in time $O((m + n)\alpha(m + n, n))$, where α is an inverse of the Ackermann function.

Initial Results

Theorem

We can solve the Offline Least Common Ancestor Problem in time $O((m+n)\alpha(m+n, n))$, where α is an inverse of the Ackermann function.

Here was Tarjan's original algorithm:

algorithm. Let T be a tree with root r and let $pairs = \{\{v_i, w_i\} \mid 1 \leq i \leq m\}$ be a set of m vertex pairs. We wish to compute $LCA(v_i, w_i)$ for each pair. The following algorithm carries out the computation.

```
procedure LCA;  
  begin  
    for each  $\{v, w\} \in pairs$  do unmark  $\{v, w\}$  od;  
    for each  $v \in V$  do create a set  $\{v\}$  named  $v$  od;  
    SEARCH( $r$ )  
  end LCA;
```

Recursive procedure SEARCH is defined by

```
procedure SEARCH( $v$ );  
  begin  
    for each  $w \in children(v)$  do SEARCH( $w$ ); UNION( $v, w$ ) od;  
    for each  $\{v, w\} \in pairs$  do if  $\{v, w\}$  not marked then mark  $\{v, w\}$   
      else  $lca(v, w) := FIND(w)$  fi od  
  end SEARCH;
```

What on earth are UNION and FIND?

[FORCED DISCUSSION FRAME]

The Union-Find Data Structure

Union-Find is a data structure that provides operations on a family of disjoint sets:

The Union-Find Data Structure

Union-Find is a data structure that provides operations on a family of disjoint sets:

- **CREATE(x)**: adds $\{x\}$ to the data structure.

The Union-Find Data Structure

Union-Find is a data structure that provides operations on a family of disjoint sets:

- $\text{CREATE}(x)$: adds $\{x\}$ to the data structure.
- $\text{UNION}(x, y)$: merges the sets in the data structure containing x and y .

The Union-Find Data Structure

Union-Find is a data structure that provides operations on a family of disjoint sets:

- $\text{CREATE}(x)$: adds $\{x\}$ to the data structure.
- $\text{UNION}(x, y)$: merges the sets in the data structure containing x and y .
- $\text{FIND}(x)$: returns a fixed representative element in the set in the data structure containing x .

Example

CREATE(x):

$\{x\}$

Example

CREATE(x):

$\{x\}$

CREATE(y):

$\{x\} \quad \{y\}$

Example

CREATE(x):

$\{x\}$

CREATE(y):

$\{x\} \quad \{y\}$

CREATE(z):

$\{x\} \quad \{y\} \quad \{z\}$

Example

CREATE(x):

$\{\mathbf{x}\}$

CREATE(y):

$\{\mathbf{x}\} \quad \{\mathbf{y}\}$

CREATE(z):

$\{\mathbf{x}\} \quad \{\mathbf{y}\} \quad \{\mathbf{z}\}$

UNION(x, y):

$\{x, \mathbf{y}\} \quad \{\mathbf{z}\}$

Example

CREATE(x):

$\{\mathbf{x}\}$

CREATE(y):

$\{\mathbf{x}\} \quad \{\mathbf{y}\}$

CREATE(z):

$\{\mathbf{x}\} \quad \{\mathbf{y}\} \quad \{\mathbf{z}\}$

UNION(x, y):

$\{x, \mathbf{y}\} \quad \{\mathbf{z}\}$

UNION(x, z):

$\{\mathbf{x}, y, z\}$

[PAUSE FRAME]

Tarjan's Union-Find I

Tarjan represents each set as a tree; the root serves as the representative element.

Tarjan's Union-Find I

Tarjan represents each set as a tree; the root serves as the representative element.




FIG 1. A *FIND* on element a , with collapsing. Triangles denote subtrees. Collapsing converts tree T into tree T' .




FIG 2 Union of two trees. Root r_1 of T_1 has a descendants; root r_2 of T_2 has b descendants. Root of new tree has $a + b$ descendants.

Tarjan's Union-Find II

More formally, those pictures become the following:

Tarjan's Union-Find II

More formally, those pictures become the following:

Find(x)

if $x.\text{parent} \neq x$

$x.\text{parent} \leftarrow \text{Find}(x.\text{parent})$

return $x.\text{parent}$

Tarjan's Union-Find II

More formally, those pictures become the following:

Find(x)

```
if x.parent  $\neq$  x
    x.parent  $\leftarrow$  Find(x.parent)
return x.parent
```

Union(x, y)

```
if x = y
    return
if x.rank < y.rank
    x.parent  $\leftarrow$  y
else if x.rank > y.rank
    y.parent  $\leftarrow$  x
else
    y.parent  $\leftarrow$  x
    x.rank  $\leftarrow$  x.rank + 1
```


Tarjan's Union-Find II

More formally, those pictures become the following:

Find(x)

```
if x.parent  $\neq$  x
    x.parent  $\leftarrow$  Find(x.parent)
return x.parent
```

Union(x, y)

```
x  $\leftarrow$  Find(x)
y  $\leftarrow$  Find(y)
if x = y
    return
if x.rank < y.rank
    x.parent  $\leftarrow$  y
else if x.rank > y.rank
    y.parent  $\leftarrow$  x
else
    y.parent  $\leftarrow$  x
    x.rank  $\leftarrow$  x.rank + 1
```

Tarjan's Union-Find III

This data structure performs well:

Tarjan's Union-Find III

This data structure performs well:

Theorem

Given a family of disjoint sets partitioning n elements, a sequence of $m \geq n$ FINDs and $n - 1$ UNIONS take time $\Theta(m\alpha(m, n))$, where α is an inverse of the Ackermann function.

Tarjan's Union-Find III

This data structure performs well:

Theorem

Given a family of disjoint sets partitioning n elements, a sequence of $m \geq n$ FINDs and $n - 1$ UNIONS take time $\Theta(m\alpha(m, n))$, where α is an inverse of the Ackermann function.

Proof.

Eight dense pages of **black magic**. □

[PAUSE FRAME]

Getting Back On-Track

Here was Tarjan's original algorithm:

Getting Back On-Track

Here was Tarjan's original algorithm:

algorithm. Let T be a tree with root r and let $pairs = \{\{v_i, w_i\} \mid 1 \leq i \leq m\}$ be a set of m vertex pairs. We wish to compute $LCA(v_i, w_i)$ for each pair. The following algorithm carries out the computation.

```
procedure LCA;  
begin  
  for each  $\{v, w\} \in pairs$  do unmark  $\{v, w\}$  od;  
  for each  $v \in V$  do create a set  $\{v\}$  named  $v$  od;  
  SEARCH( $r$ )  
end LCA;
```

Recursive procedure SEARCH is defined by

```
procedure SEARCH( $v$ );  
begin  
  for each  $w \in children(v)$  do SEARCH( $w$ ); UNION( $v, w$ ) od;  
  for each  $\{v, w\} \in pairs$  do if  $\{v, w\}$  not marked then mark  $\{v, w\}$   
    else  $lca(v, w) := FIND(w)$  fi od  
end SEARCH;
```

Getting Back On-Track

Here was Tarjan's original algorithm:

algorithm. Let T be a tree with root r and let $pairs = \{\{v_i, w_i\} \mid 1 \leq i \leq m\}$ be a set of m vertex pairs. We wish to compute $LCA(v_i, w_i)$ for each pair. The following algorithm carries out the computation.

```
procedure LCA;  
begin  
  for each  $\{v, w\} \in pairs$  do unmark  $\{v, w\}$  od;  
  for each  $v \in V$  do create a set  $\{v\}$  named  $v$  od;  
  SEARCH( $r$ )  
end LCA;
```

Recursive procedure SEARCH is defined by

```
procedure SEARCH( $v$ );  
begin  
  for each  $w \in children(v)$  do SEARCH( $w$ ); UNION( $v, w$ ) od;  
  for each  $\{v, w\} \in pairs$  do if  $\{v, w\}$  not marked then mark  $\{v, w\}$   
    else  $lca(v, w) := FIND(w)$  fi od  
end SEARCH;
```

Let's parse this with an example.

Example


Pairs:

- b, d
- f, e

Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:


- b, d
- f, e



Example

Pairs:

- b, d
- f, e



[PAUSE FRAME]

TarjanOLCA(u)

```
u.ancestor ← u
for v ∈ u.children do
    TarjanOLCA(v)
    Union(u,v)
    Find(u).ancestor ← u
u.colour ← black
for v such that {u,v} ∈ pairs do
    if v.colour = black
        Least Common Ancestor of u and v is Find(v).ancestor
```

[FORCED DISCUSSION FRAME]

Analyzing the Code

Now that we finally understand the code, let's analyze its runtime:

Analyzing the Code

Now that we finally understand the code, let's analyze its runtime:

- The first loop implies each edge causes one `UNION` and `FIND`.

Analyzing the Code

Now that we finally understand the code, let's analyze its runtime:

- The first loop implies each edge causes one `UNION` and `FIND`.
- The second loop implies each pair causes one `FIND`.

Analyzing the Code

Now that we finally understand the code, let's analyze its runtime:

- The first loop implies each edge causes one UNION and FIND.
- The second loop implies each pair causes one FIND.
- Thus this algorithm invokes $m + n - 1$ FINDs and $n - 1$ UNIONS.

Analyzing the Code

Now that we finally understand the code, let's analyze its runtime:

- The first loop implies each edge causes one UNION and FIND.
- The second loop implies each pair causes one FIND.
- Thus this algorithm invokes $m + n - 1$ FINDs and $n - 1$ UNIONS.

Theorem

We can solve the Offline Least Common Ancestor Problem in time $O((m + n)\alpha(m + n, n))$, where α is an inverse of the Ackermann function.

[END FRAME]